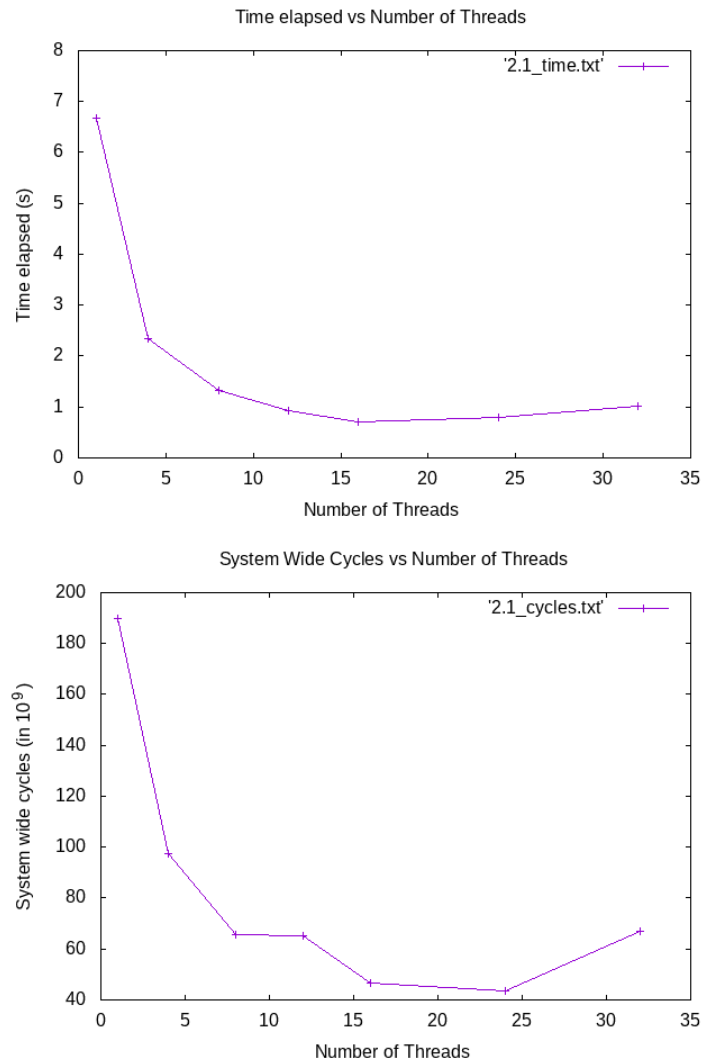


1 Perf Stat

Note: All readings here are done for 6 runs of the code, i. e. the default number of runs.



The pattern observed is that both system-wide cycles and time elapsed decrease when the number of threads increase from 1 but start increasing when the number crosses a certain limit (Which appears to be 16 threads here).

The reason for this is that initially, the threads are able to effectively use the parallelism available on the cluster machine by assigning each thread to a different core/CPU which reduces the overall runtime and hence the number of cycles. However, when the number of threads becomes too high, each thread cannot be assigned to a different core/CPU so the whole task cannot be done parallelly.

2 Perf Record

The assembly instruction which takes the most time (37.95%) is **lg 93**.

Using the **-g** compile flag in the **CFLAGS** variable in the makefile, we can see the source code along with the assembly which allows us to determine the source code portion this instruction maps to.

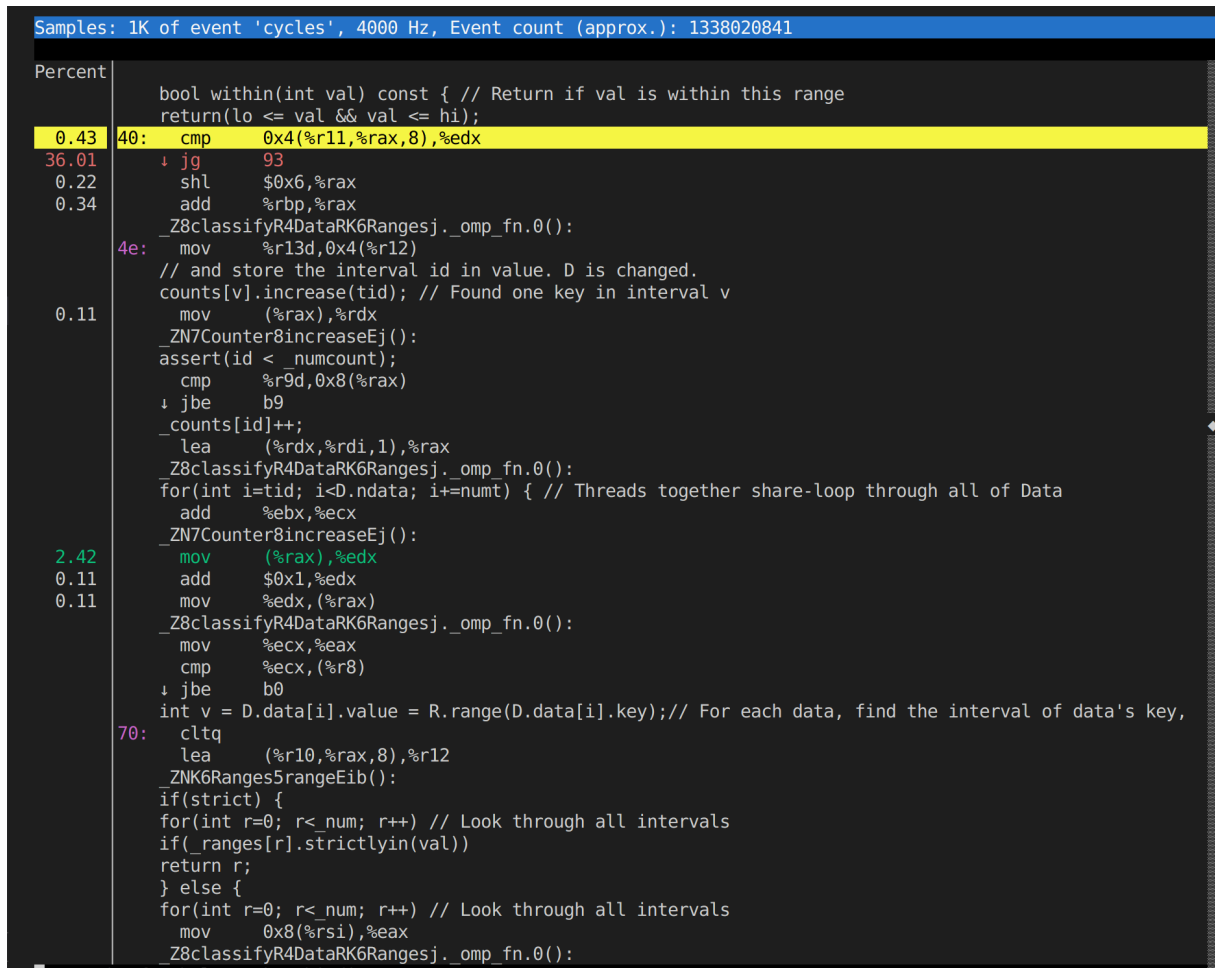
The part of the source code which maps to this instruction is:

```

if(_ranges[r].within(val))
    return r;
bool within(int val) const { // Return if val is within this range
    return(lo <= val && val <= hi);
}

```

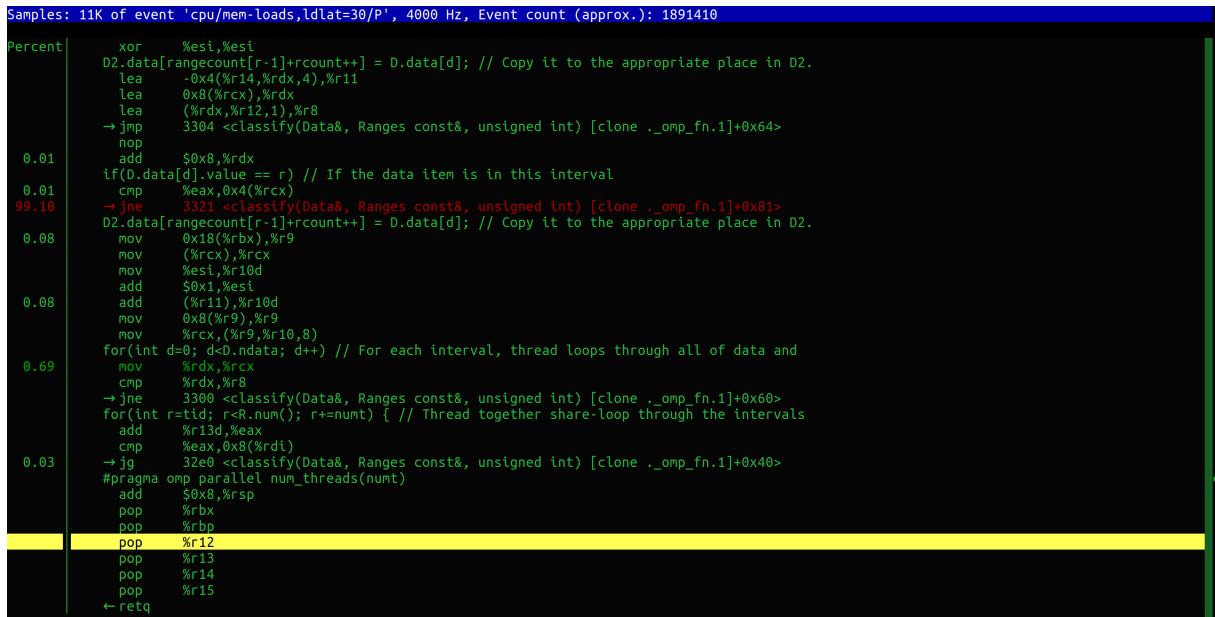
3 Hotspot Analysis



The problem which makes this portion of code the bottleneck lies in the classification algorithm. The classification algorithm uses linear search to determine the range a value lies in which causes this function to be called much more than it needs to.

The code can be optimized by changing the linear search to a binary search.

4 Memory Profiling



From the above image we can see that the top 2 hotspots in the code are **jne 3321** (99.10% of time) and **mov %rdx,%rcx** (0.69% of time).

The instruction **jne 3321** corresponds to:

```
for(int d=0; d<D.ndata; d++) // For each interval, thread loops through all of data and
    if(D.data[d].value == r) // If the data item is in this interval
        D2.data[rangeount[r-1]+rcount++] = D.data[d]; // Copy it to the appropriate place in D2.
```

The instruction **mov %rdx,%rcx** corresponds to:

```
for(int d=0; d<D.ndata; d++) // For each interval, thread loops through all of data and
```

In the portion of the code shown in the first hotspot, each thread works on the data points which have modulo of the thread ID. This is cache unfriendly because different threads are attempting to access contiguous portions of memory, i. e. Thread 0 accesses address 0, thread 1 accesses address 1, etc. which are likely on the same cache line which leads to the threads accessing the same cache line concurrently leading to false sharing.

This problem can be made cache friendly by having each thread itself do a contiguous portion of the work, i. e. for 100 samples and 4 threads, thread 0 does 0-24 and so on. This leads to minimization of concurrent same cache line accesses by different threads.

On observing the code, we see another case of false sharing when the interval of the data's key is found, having the same issue as the previous case and therefore the same solution.

On optimizing we get:

Percent	<pre> nop int rcount = 0; for(int d=0; d<D.ndata; d++) // For each interval, thread loops through all of data and mov (%rbx),%rdx mov (%rdx),%eax test %eax,%eax → je 3343 <classify(Data&, Ranges const&, unsigned int) [clone ._omp_fn.1]+0xa3> mov 0x8(%rdx),%rdx lea -0x1(%rax),%esi lea 0x8(%rdx),%rax lea (%rax,%rsi,8),%r8 int rcount = 0; xor %esi,%esi → jmp 331c <classify(Data&, Ranges const&, unsigned int) [clone ._omp_fn.1]+0x7c> nop add \$0x8,%rax if(D.data[d].value == r) // If the data item is in this interval 48.22 cmp 0x4(%rdx),%ecx → jne 333b <classify(Data&, Ranges const&, unsigned int) [clone ._omp_fn.1]+0x9b> 1.55 D2.data[rangecount[r-1]+rcount++] = D.data[d]; // Copy it to the appropriate place in D2. 49.37 mov 0x18(%rbx),%r10 mov (%rdx),%rdx mov %esi,%r11d add \$0x1,%esi 0.46 add -0x4(%rbp,%r9,4),%r11d 0.39 mov 0x8(%r10),%r10 mov %rdx,(%r10,%r11,8) </pre>
	<pre> for(int d=0; d<D.ndata; d++) // For each interval, thread loops through all of data and mov %rax,%rdx cmp %rax,%r8 → jne 3318 <classify(Data&, Ranges const&, unsigned int) [clone ._omp_fn.1]+0x78> for(int r=start; r<start+((tid==numt-1)?div+mod:div); r++) { // Thread together share-loop through th add \$0x1,%r9 mov %r9d,%ecx cmp %r9d,%edi → jg 32f8 <classify(Data&, Ranges const&, unsigned int) [clone ._omp_fn.1]+0x58> #pragma omp parallel num_threads(numt) pop %rbx pop %rbp pop %r12 pop %r13 </pre>

Initial cache misses: 4746131

Final cache misses: 4177119

We see an improvement in cache misses.