

1 Common Approach

The algorithm used to find the truss value of each edge is the same for both task 1 and task 2. The algorithm implemented is the hybrid truss algorithm. A directed version of the graph is stored at each rank, and the triangles are enumerated for a subset of edges.

Through inter-process communication, the values of each edge reduce monotonically and eventually converge from the support of the edge to the final truss value.

2 Task 1

2.1 Approach


Not Verbose:

Check if a truss of the given size exists.

Verbose:

If a truss exists, get all existing edges and run BFS to find connected components.

2.2 Speedup & Efficiency



2.3 Iso-Efficiency

After observing the results, the best function for this algorithm's iso efficiency is $\Omega(p^{1.5})$.

2.4 Sequential Fraction Estimation

The sequential fraction of the algorithm is strictly decreasing as the number of processes and threads increases. Therefore I estimate that the sequential fraction of the algorithm will be approximately 0.15, 0.1, 0.08 for 16, 32, 64 cores, respectively.

2.5 Scalability

The algorithm depends on the graph structure, so drawing a concrete conclusion about the scalability is difficult. However, from the figures, it is visible that the problem is scaling decently as cores increase.

3 Task 2

3.1 Approach

Take the connected components of the k-truss - found using BFS.

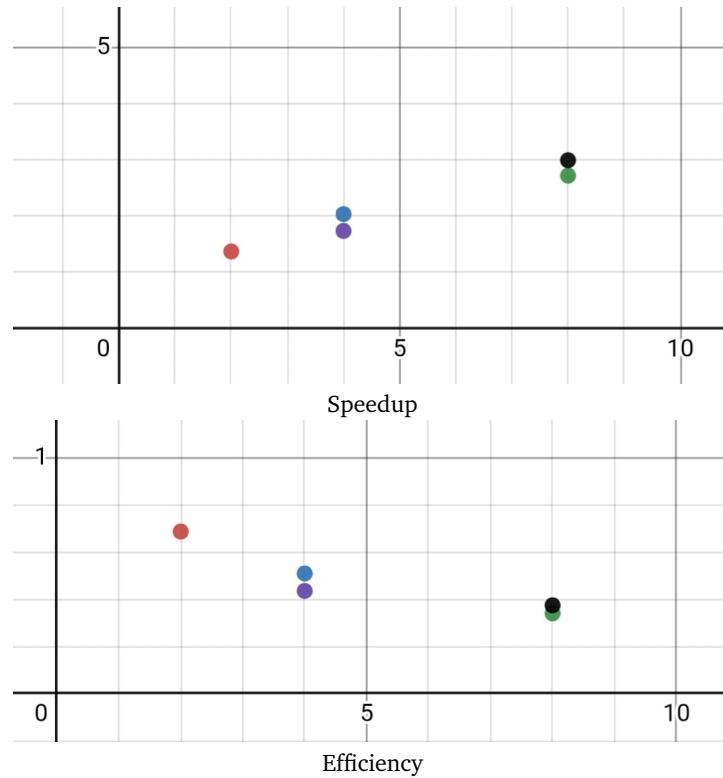
Not Verbose

Check which vertices are connected to a sufficient number of components.

Verbose

Find the vertices connected to a sufficient number of components and output those components.

3.2 Speedup & Efficiency



3.3 Iso-Efficiency

After observing the results, the best function for this algorithm's iso efficiency is $\Omega(p^{1.5})$.

3.4 Sequential Fraction Estimation

The sequential fraction of the algorithm is strictly decreasing as the number of processes and threads increases. Therefore I estimate that the sequential fraction of the algorithm will be approximately 0.15, 0.1, 0.08 for 16, 32, 64 cores, respectively.

3.5 Scalability

The algorithm depends on the graph structure, so drawing a concrete conclusion about the scalability is difficult. However, from the figures, it is visible that the problem is scaling decently as cores increase.